

NetApp, Inc. (NASDAQ: NTAP)

Market Cap: \$20.15 B
#1 spot in hybrid storage

Intelligent data infrastructure that unifies storage across flash, hybrid, and cloud

Company Overview

- NetApp is a San Jose-based data infrastructure company founded in 1992 that provides enterprise storage, data management, and hybrid cloud solutions to large organizations.
- NetApp has two main product categories: hybrid cloud, and public cloud.
 - Hybrid cloud includes businesses' own isolated data center infrastructure (private server/on-prem) and collaborating with public clouds to achieve control, security and flexibility.
 - For on-prem physical hardware, FAS Software offers perpetual licensing fees, ranging from \$20,000 to several hundred thousand dollars for larger enterprises
 - FAS Software installed on private server, licensed on a capacity-based model, has costs ranging roughly from \$0.04-\$0.1 per GB per year for licenses (ONTAP manages 100+ EB of data)
 - NetApp also provides first-party services (NAS) by integrating its ONTAP storage operating system directly into AWS, Azure, and Google Cloud.
 - NAS charges subscription fee to public clouds with prices often starting around \$0.03-\$0.04 per GB per month. Depending on the size of customer, revenue could range from 10TB to several PB.
- Management has a history of weak capital allocation, failed adjacencies, and internal instability
 - Under George Kurian, NetApp spent about \$1.3B on cloud acquisitions over five years, including more than \$450M for Spot.io, CloudCheckr, and Data Mechanics, then sold the Spot and CloudCheckr portfolio to Flexera in January 2025 for roughly \$100M.

Overview & Trends

- The global datasphere is expanding rapidly, which drives storage demand.
 - IDC's Global DataSphere forecast implies global data rises from ~213.6 zettabytes in 2025 to ~527.5 zettabytes in 2029 (~ 2.5x in four years), which will lead to data storage demand increase.
 - The marginal driver is AI: more model training/fine-tuning, more telemetry, and more unstructured content (audio/video/image) all increase data creation and retention.
- Incorporation of AI in these platforms raise both capacity demand and performance
 - AI training and inference can impose a performance bottleneck and cost of moving data from storage to GPUs is high. This has supported the rise of AI-optimized storage, including parallel file systems that prioritize speed, metadata performance, and concurrency over traditional enterprise storage design.
 - Specialist vendors are taking share in pockets where requirements are extreme, especially AI pipelines that demand throughput and high-concurrency access, but traditional vendors also experience growth.
- Hybrid cloud demand should increase as enterprises demand speed and economics
 - Regulated / sensitive datasets often face privacy, residency, and governance constraints, keeping meaningful data on-prem or in private environments, while less sensitive workloads shift to public clouds for agility and cost optimization.
 - 90% of organizations will use hybrid cloud structure through 2028, reinforcing that "hybrid" is becoming the default operating model rather than a transition phase.
- Transmission become friction points, making "bridges" like ONTAP strategically important but competitive
 - Gartner flags data synchronization across the hybrid cloud environment as an urgent GenAI challenge, which elevates value of storage/data layers that can move and manage data across environments reliably.
 - Storage environments are sticky as switching costs are high. Customers prioritize data integrity, ransomware resilience, uptime, and operational continuity, which increases inertia once a platform is embedded in workflows.

Competitors Separated by Service Type

1. Private Storage Arrays

- Storage arrays are the physical hardware companies use to hold data in their own data centers. Historically, the buying decision centered on how much data could be consolidated under a single management layer with consistent policies and low operational overhead. NetApp built its position around that: ONTAP, its operating system, handles multiple data types with integrated backup, security, and policy enforcement. But AI workloads are shifting the criteria. AI infrastructure needs to move massive volumes of data at high speed to keep GPUs utilized, and procurement increasingly sits with ML engineering teams rather than traditional IT. Those buyers optimize for raw throughput per hardware unit and cost per unit of speed, not the breadth of bundled management features. NetApp's latest system requires roughly twice the hardware footprint of Pure Storage to deliver comparable throughput, because ONTAP's enterprise feature set creates overhead at scale. As AI storage budgets grow faster than general-purpose enterprise storage, that performance gap becomes an increasingly material competitive disadvantage.

2. Bridge Systems / Data Fabric Vendors

- Most large enterprises operate a mix of private data centers and public cloud infrastructure from AWS, Microsoft Azure, or Google Cloud. Maintaining consistent data protection, security, and compliance across those environments is operationally complex, and NetApp's core value proposition is that ONTAP provides a unified set of governance, replication, and recovery controls regardless of where data sits. That advantage is genuine, but it is narrower than it appears. It is most relevant for regulated, compliance-driven workloads where auditability and consistent policy enforcement are hard requirements. For AI workloads, where the fastest-growing storage spending is concentrating, engineering teams prioritize throughput and operational simplicity over cross-environment governance. Pure Storage offers a leaner, higher-performance model that is increasingly competitive for those buyers. NetApp's governance strength is real but risks becoming a niche advantage while the highest-value workloads move to purpose-built competitors.

3. Collaboration with Public Clouds

- NetApp has a structurally distinct arrangement: its storage software runs as a managed service inside the major clouds, with Amazon, Microsoft, and Google handling sales, provisioning, and support on NetApp's behalf. That provides broad distribution, but it also makes NetApp a dependent layer inside platforms it does not control. The cloud providers set the pricing, own the customer relationship, and determine the product roadmap. They are also steadily building native versions of the data management capabilities that differentiate NetApp's software. As those tools approach good-enough parity, particularly for AI workloads where customers want the most direct path to high-throughput storage, the justification for an additional third-party layer weakens. Pure Storage takes a different approach, extending its own system into the cloud while retaining direct control of customer relationships and pricing, which may prove more defensible if cloud providers move from partnering with third-party storage vendors to competing with them.

Amy Li, Jiong Li,
Jonathan Plavnik, James Wang

March 14, 2026

Key Ratios	
Recommendation	Sell
Consensus	Buy
Share Price	\$98.66
Price Target	\$64.97
52-Week Low	\$71.8
52-Week High	\$127.8
Net Debt/EBITDA	1.6x
EV/Sales	2.8x
EV/EBITDA	10.6x
Price/Book	18.2x



Metrics				
\$M	FY 2022	FY 2023	FY 2024	FY 2025
Revenue	6,318.0	6,362.0	6,268.0	6,572.0
EBITDA	1,460.0	1,468.0	1,571.0	1,709.0
Net Debt	2,940.0	2,684.0	2,652.0	3,491.0
EPS, Adj	4.26	6.30	4.82	5.77

Revenue Breakdown				
\$M	FY 2022	FY 2023	FY 2024	FY 2025
Hybrid Cloud	5,922.0	5,787.0	5,657.0	5,907.0
Product	3,284.0	3,049.0	2,849.0	3,040.0
Hardware	1,358.0	1,251.0	1,251.0	No Data
Software	1,926.0	1,798.0	1,598.0	No Data
Support	2,638.0	2,738.0	2,808.0	2,867.0
Public Cloud	396.0	575.0	611.0	665.0
Total Revenue	6,318.0	6,362.0	6,268.0	6,572.0

Trading	
Average Volume	2.22M
50-Day Moving Average	107.91M
200-Day Moving Average	107.92M

Risk	
Beta	1.39
Insider Ownership	0.50%
Short % of Float	9.49%

Price Chart & Scenario Analysis	
1. Base Case Fair Value: \$64.97 (- 36.4%)	
2. Bull Case Fair Value: \$54.71 (- 46.4%)	
3. Bear Case Fair Value: \$100.8 (- 1.3%)	



Investment Thesis

1. NetApp's gross margin is more vulnerable than the market appreciates because the company still derives a substantial portion of its economics from hardware-linked storage systems, which leaves product margins exposed to NAND and broader component inflation.

- Investors increasingly view NetApp as an AI-enabled data infrastructure platform, but the underlying margin profile is still anchored by the Hybrid Cloud segment, where product gross margins remain sensitive to changes in hardware, memory, controller, and networking input costs. NetApp already disclosed that Hybrid Cloud product gross margins fell by roughly two points in FY2025 due primarily to higher component costs, and Q3 FY2026 product gross margin fell again to 55.3% in Q3 FY2026 versus 59.6% consensus. The market treats these pressures as temporary, but industry supply data suggest a broader cost stack, particularly server DRAM and high-capacity disk supply, which may remain tight into 2027, making a clean margin rebound harder than consensus assumes.
- Even if costs remain elevated, NetApp may not be able to fully protect margin because the storage market is highly competitive. The issue is not just inflation, but inflation in a market where customers can compare across Dell, HPE, Pure, and hyperscaler-native options. NetApp itself says gross margins are affected by competitive pricing, discounting pressure, and the evolving competitive landscape shaped by public cloud providers. That makes full cost pass-through harder than the market may be assuming.
- The bull case leans heavily on the shift toward software and cloud, but that transition may take longer to show up in margins than investors expect. Public Cloud is growing, but it is still only 10% of total revenue, and scaling it requires continued spending on sales and hyperscaler partnerships. At the same time, the support and services business that gives NetApp its margin stability is ultimately tied to how much hardware is in the field. Fewer product shipments today means a smaller renewal base down the road, which erodes the recurring revenue stream the market is counting on.

2. NetApp's enterprise DNA is a liability in AI infrastructure, as performance-per-node and architectural simplicity are the buying criteria for most customers.

- Enterprise AI is moving from pilot to production. About 80% of Fortune 500 companies are experimenting with GenAI, but less than 20% have production deployments at scale. When this flips, likely around Q2 2026, storage demand shifts to speed and capacity. Pilot workloads can tolerate whatever infrastructure is already in place, but large-scale deployments have much higher requirements. AI workloads create massive unstructured datasets with high latency sensitivity, and that transition from pilot to production is exactly when teams re-evaluate their storage stack. NetApp's largest customers are also the ones moving most aggressively to build AI infrastructure, so as budgets shift away from general-purpose storage and toward performance-optimized systems, NetApp stands to lose meaningful share in its most valuable accounts. That is why the bull case that NetApp simply holds its position because it already sits inside those enterprise accounts is unlikely to happen.
- NetApp is picking up some early AI-adjacent storage deals, but it has a structural handicap. ONTAP's enterprise feature set adds overhead that forces heavier scale-out to match purpose-built competitors. NetApp's AFX can scale to 4 TB/s of read throughput, but needs up to 128 controllers, roughly twice the hardware footprint that competitive parallel storage systems require for comparable bandwidth. AI buyers optimize for GPU utilization and \$/GB/s, not unified protocol support, so a system that needs 2x the nodes to hit the same throughput is a harder sell when the procurement decision sits with ML engineering teams rather than traditional IT.
- NetApp's pivot toward AI is slow. It is layering AI positioning onto an architecture built for general-purpose enterprise storage, not rebuilding from scratch. That means it can capture the governance and hybrid layer around AI workflows but is structurally weaker for the highest-throughput training and inference workloads where the largest incremental storage dollars are concentrating. The risk is that "AI-adjacent" becomes a ceiling: NetApp services the compliance and data management perimeter while competitors like Pure Storage and emerging parallel file systems take the performance core, which is where the premium multiples and fastest growth sit.

3. NetApp's core storage architecture is structurally more vulnerable than the market appreciates, as disaggregated platforms, hyperscaler-native storage, and more efficient next-generation hardware each attack a different layer of the legacy ONTAP-led model.

- NetApp's software layer is more vulnerable than the market appreciates because ONTAP is increasingly competing in a part of the stack where differentiation is difficult and legacy complexity becomes a disadvantage. As more of NetApp's earnings base is underwritten by software and recurring services rather than proprietary hardware, the company sounds more financially durable on the surface, but it also means NetApp is moving into areas that both hyperscalers and next-generation storage vendors can build natively without carrying ONTAP's architectural baggage. In June 2025, former SVP and CTO Jonsi Stefánsson left NetApp, formed Red Stapler with other former core NetApp engineers, and sold it to VAST Data a few months later. NetApp sued, alleging he used company IP to build a competing platform, but the case was dismissed in Florida in February 2026, with NetApp appealing afterward. Regardless of the legal outcome, the takeaway is that one of the executives closest to NetApp's cloud architecture chose to build the next platform outside the company rather than inside it. Vendors like VAST are not trying to beat NetApp at legacy enterprise storage; they are bypassing it with cleaner software environments designed around scale-out flash architectures and lower operating complexity. The key point is that ONTAP's installed base matters less when customers are evaluating software around workload-specific optimization.
- NetApp's Public Cloud is delivered as-a-service with little upfront purchase, so hyperscalers (AWS, Azure, and GCP) pay as they consume capacity and NetApp recognizes revenue as usage occurs, producing smoother revenue, higher gross margins, and greater stickiness because customer workflows are directly on NetApp's platform rather than purchased hardware. But that same model leaves NetApp increasingly at the mercy of hyperscalers, which own the marketplace, control pricing and distribution, and can compress NetApp's economics by launching competing products or changing commercial terms. The durability investors see in the model comes with a loss of strategic control: NetApp no longer fully owns the customer relationship and is instead relying on infrastructure controlled by the very platforms most likely to take over its function.
- NetApp's legacy hardware business faces pressure from more efficient, specialized players like Pure Storage, and hybrid persistence does not solve that problem because it keeps on-prem demand intact without guaranteeing NetApp wins the hardware refresh. At scale, NetApp deployments can require as many as 128 nodes to achieve performance that competitors like Pure can deliver with fewer than half, which drives a materially worse total cost of ownership. At the same time, AWS Outposts, Azure Stack, and Google Distributed Cloud reduce the value of NetApp's bridge position by embedding hybrid extensions directly into the hyperscaler stack, leaving NetApp's hardware business increasingly exposed as customers reassess what infrastructure they actually need on-prem. Bulls point to the 54% of revenue from support and services as a durable subscription-like layer, but that stream is still tied to the installed hardware base. If product shipments slow because customers choose denser, more efficient alternatives, the renewal pool shrinks with it.

Catalysts

- Enterprise AI moves from pilot to production at scale. As Fortune 500 companies shift from experimentation to full deployment, likely within two quarters, storage procurement moves from general-purpose to performance-specific budgets. That transition triggers re-evaluation inside NetApp's most valuable customers. The bull case assumes NetApp holds those accounts by default, but the more likely outcome is that production AI workloads get purpose-built from scratch, and the procurement process favors vendors who win on throughput per node and cost per GB/s. Thus, every large account that re-platforms its AI storage stack is a permanent share loss.
- Hyperscalers improve and bundle native storage capabilities. AWS, Azure, and GCP are steadily building out the data tiering, orchestration, and lifecycle management features that NetApp sells as its software differentiation. When those features become good enough inside native cloud storage and get bundled into enterprise discount programs, the value of running NetApp as a dependent layer inside those marketplaces compresses. NetApp's public cloud revenue (\$665M, 10% of total) is presented as a growth engine, but it is entirely dependent on hyperscalers continuing to promote a third-party solution rather than competing with it. Any shift in that dynamic, whether through first-party feature parity, repricing of partner economics, or preferential bundling, directly pressures the revenue stream the bull case relies on for a re-rate.
- Regulatory and data-sovereignty requirements intensify. This is a tailwind for the bull case, but the same regulatory pressure also benefits hyperscalers on-prem extensions like AWS Outposts, Azure Stack, and Google Distributed Cloud, which deliver compliant private environments with a first-party cloud experience. As those offerings mature, enterprises can satisfy governance requirements without a third-party storage vendor in the stack, making NetApp's bridge role substitutable rather than reinforced.

Risks & Mitigants

- NetApp's installed-base stickiness proves more durable than expected. If ONTAP is so deeply embedded in enterprise recovery, governance, and replication workflows that rip-and-replace costs remain high even as AI budgets grow, NetApp could retain its existing revenue base while adding AI-adjacent workloads at the margin. If customers treat ONTAP as the model for protection and compliance while adding separate AI storage alongside it rather than replacing it, NetApp avoids too much share loss.
- NetApp's AFX or a future architecture iteration closes the performance gap. If NetApp can materially reduce the node count required to hit competitive throughput levels, the structural disadvantage in AI workloads narrows. The current 2x hardware footprint gap is based on early AFX feedback, and subsequent generations could improve scale-out efficiency enough that engineering teams consider NetApp viable for production AI, particularly if it is bundled with the governance and hybrid capabilities pure-play competitors don't possess.
- Hyperscalers continue to promote NetApp rather than compete with it. NetApp's cloud revenue depends on AWS, Azure, and GCP treating ONTAP as the recommended enterprise storage layer. If that partnership dynamic holds because hyperscalers prefer to focus on computing, networking, and AI services rather than building out enterprise-grade storage management themselves, then NetApp's consumption-based cloud revenue could continue growing.
- A recession delays enterprise AI deployment timelines. If macro conditions push production AI deployments further out, NetApp's largest customers stay on general-purpose storage longer, deferring the vendor re-evaluation and giving management more time to improve the product and close the performance gap.